

Rex Chao

📞 858-408-5697 | 📩 rexchao.careers@gmail.com | 💼 linkedin.com/in/rexchao | 🐾 github.com/seraph27 | 🌐 seraph.tw

EDUCATION

University of California San Diego <i>B.S. in Computer Science and Mathematics / Major GPA: 3.9/4.0</i>	Sep 2021 – Dec 2025 (Expected) La Jolla, CA
• Coursework: Design & Analysis of Algorithms · Recommender Systems & Web Mining · Machine Learning Algorithms · Advanced Data Structures · Data Science in Practice · Intro to Stochastic Processes · Intro to Numerical Analysis	

EXPERIENCE

Lunabill — Y Combinator F25 <i>Founding Software Engineer</i>	Sep 2025 – Present San Francisco, CA
• Built real-time observability dashboard processing 5,000+ calls/wk across \$25M+ in insurance claims, automating call status tracking and re-dial workflows that eliminated 200+ hrs/mo of manual rescheduling and triage work.	
Foxconn iAi <i>Machine Learning Engineer Intern</i>	Jul 2025 – Sep 2025 Milwaukee, WI
• Optimized real-time factory voice-agent pipeline (Whisper ASR + LLaMA + RAG) enabling technicians to query maintenance history and equipment logs via natural language instead of manual searches.	
Formosa Dental Clinic <i>Software Engineer</i>	May 2025 – Jul 2025 Taipei, Taiwan
• Cut ASR cold-start latency by 15% via multi-threaded model loading and warm-ups and reduced end-to-end latency by 25% with streaming tweaks (sliding windows, adaptive chunking, VAD) validated by A/B tests.	
• Deployed demo in factory environment for internal testing with 12 technicians , achieving 77% query success rate and collecting feedback that validated core use cases for production consideration.	
Systex Software <i>Software Developer Intern</i>	Jul 2024 – Sep 2024 Taipei, Taiwan
• Engineered a full-stack PWA for operations—scheduling, billing, patient records, and inventory with low-stock alerts.	
• Migrated 36k+ patient records, and built a system that has processed \$1M+ in transactions .	
• Built a secure data pipeline with automated balance reconciliation and audit trails to ensure integrity and reliability.	
• Implemented automated backups (Vercel Cron + Dropbox API): daily snapshots, 30-day retention, and disaster recovery.	

PROJECTS

Tsukiyomi — AI Japanese Light-Novel Recommender Python, PyTorch, Hugging Face, SentencePiece, Milvus
• Learns a personal known-word profile from quick word marking to recommend chapters matched to a learner's level.
• Implemented dual-tower retrieval (text embeddings + user-vocab embeddings) with Milvus ANN; re-ranks via unknown-word rate, LM predictability, and genre similarity.
• Indexed over 60k+ chapters and returns top-10 prediction in 200ms .

LEADERSHIP & ACHIEVEMENTS

Officer, Competitive Programming Club @ UCSD (Jane Street Sponsored)	Feb 2024 – Present
• Organized and presented advanced algorithm workshops to over 30+ members; revamped club website.	
Competitive Programming Achievements	Present
• LeetCode: 2303 rating (Top 0.5% of 17M+ users); Codeforces: 1804 rating (Top 7% globally).	

TECHNICAL SKILLS

Stack: C++, Python, TypeScript, SQL; PyTorch, React/Next.js; PostgreSQL; Docker, GitHub Actions, Vercel, GCP, Linux.